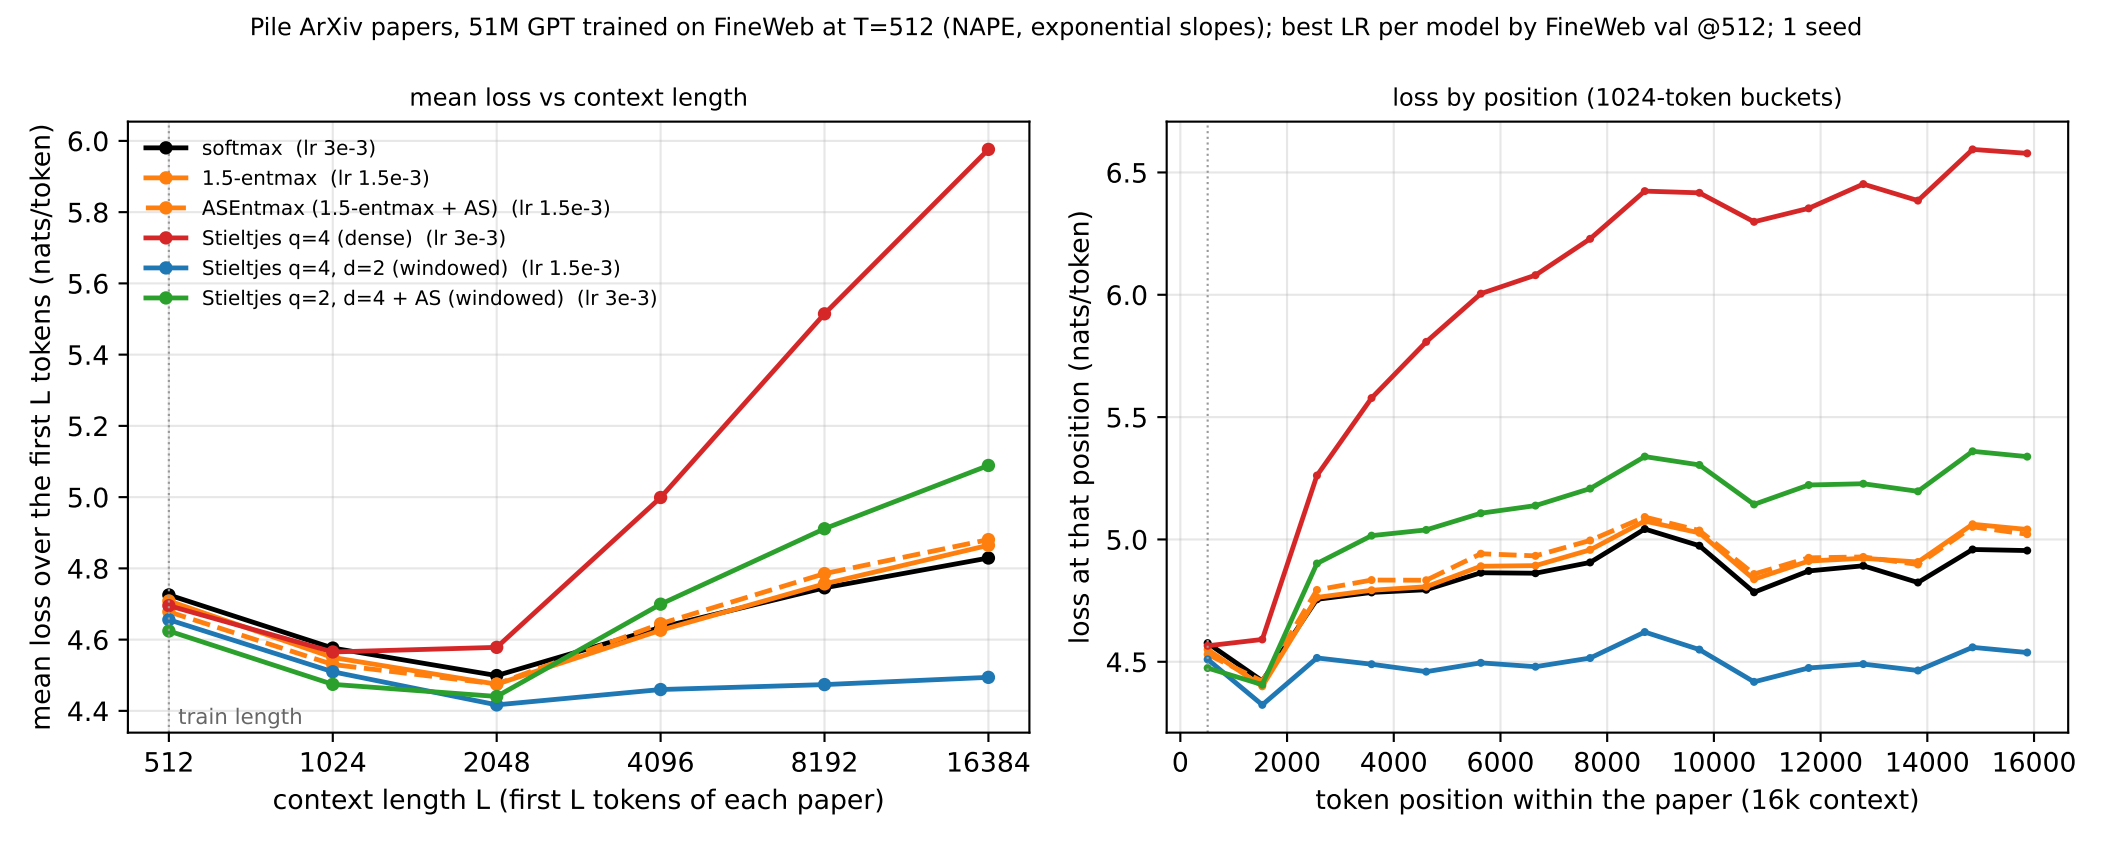
Pile ArXiv papers, 51M GPT trained on FineWeb at T=512 (NAPE, exponential slopes); best LR per model by FineWeb val @512; 1 seed
mean loss vs context length
softmax (lr 3e-3)
1.5-entmax (lr 1.5e-3)
ASEntmax (1.5-entmax + AS) (lr 1.5e-3)
Stieltjes q=4 (dense) (lr 3e-3)
Stieltjes q=4, d=2 (windowed) (lr 1.5e-3)
Stieltjes q=2, d=4 + AS (windowed) (lr 3e-3)
train length
mean loss over the first L tokens (nats/token)
context length L (first L tokens of each paper)
loss by position (1024-token buckets)
loss at that position (nats/token)
token position within the paper (16k context)