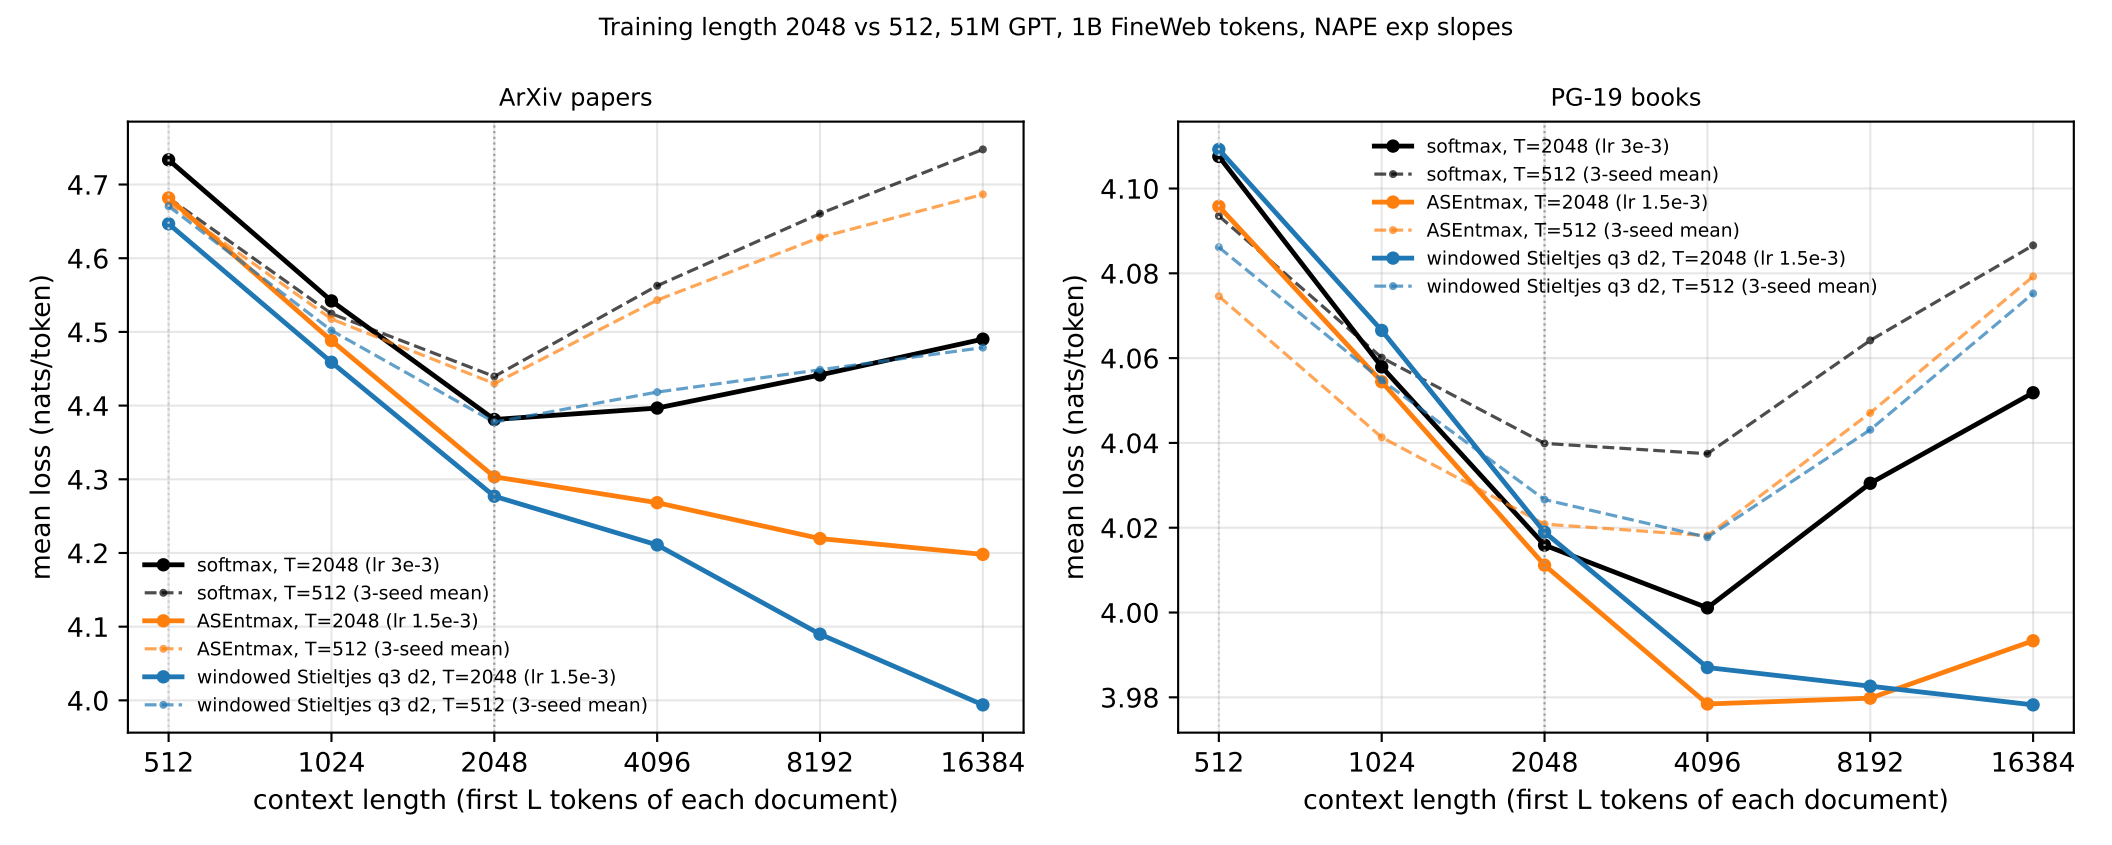
Training length 2048 vs 512, 51M GPT, 1B FineWeb tokens, NAPE exp slopes
ArXiv papers
PG-19 books
mean loss (nats/token)
context length (first L tokens of each document)
512
1024
2048
4096
8192
16384
4.0
4.1
4.2
4.3
4.4
4.5
4.6
4.7
3.98
4.00
4.02
4.04
4.06
4.08
4.10
softmax, T=2048 (lr 3e-3)
softmax, T=512 (3-seed mean)
ASEntmax, T=2048 (lr 1.5e-3)
ASEntmax, T=512 (3-seed mean)
windowed Stieltjes q3 d2, T=2048 (lr 1.5e-3)
windowed Stieltjes q3 d2, T=512 (3-seed mean)